

# **INTRO to DATA SCIENCE**

## **LECTURE 11: SUPPORT VECTOR MACHINES**

## **LAST TIME:**

- ENSEMBLE TECHNIQUES**
- PROBLEMS IN CLASSIFICATION**
- BAGGING, BOOSTING, RANDOM FORESTS**

**I. SUPPORT VECTOR MACHINES**

**II. MAXIMUM MARGIN HYPERPLANES**

**III. SLACK VARIABLES**

**IV. NONLINEAR CLASSIFICATION**

**EXERCISE:**

**V. SVM IN SCIKIT-LEARN**

# **I. SUPPORT VECTOR MACHINES**

*Q: What is a support vector machine?*

*Q: What is a support vector machine?*

*A: A binary linear classifier whose decision boundary is explicitly constructed to minimize generalization error.*

*Q: What is a support vector machine?*

*A: A binary linear classifier whose decision boundary is explicitly constructed to minimize generalization error.*

*recall:*

**binary classifier** – *solves two-class problem*

**linear classifier** – *creates linear decision boundary (in 2d)*

*Q: How is the decision boundary derived?*



*Q: How is the decision boundary derived?*

*A: Using geometric reasoning (as opposed to the algebraic reasoning we've used to derive other classifiers).*

*Q: How is the decision boundary derived?*

*A: Using geometric reasoning (as opposed to the algebraic reasoning we've used to derive other classifiers).*

**NOTE**


These are two different ways of looking at the same problem.

Familiarity with both leads to deeper understanding!

*Q: How is the decision boundary derived?*

*A: Using geometric reasoning (as opposed to the algebraic reasoning we've used to derive other classifiers).*

*The generalization error is equated with the geometric concept of margin, which is the region along the decision boundary that is free of data points.*



**FIGURE 18-4.** Two decision boundaries and their margins. Note that the vertical decision boundary has a wider margin than the other one. The arrows indicate the distance between the respective support vectors and the decision boundary.

*Q: How is the decision boundary derived?*

*A: Using **geometric reasoning** (as opposed to the algebraic reasoning we've used to derive other classifiers).*

*The goal of an SVM is to create the linear decision boundary with the largest margin. This is commonly called the **maximum margin hyperplane**.*

*Q: How is the decision boundary derived?*

*A: Using **geometric reasoning** (as opposed to the algebraic we've used to derive other classifiers).*

### NOTE

A *hyperplane* is just a high-dimensional generalization of a line.

*The goal of an SVM is to create the linear decision boundary with the largest margin. This is commonly called the **maximum margin hyperplane**.*

*Q: If SVM is a linear classifier, how can you use it for nonlinear classification?*

*Q: If SVM is a linear classifier, how can you use it for nonlinear classification?*

*A: Using a clever maneuver called the **kernel trick**.*



*Nonlinear applications of SVM rely on an implicit (nonlinear) mapping  $\Phi$  that sends vectors from the original feature space  $K$  into a higher-dimensional feature space  $K'$ .*

*Nonlinear applications of SVM rely on an implicit (nonlinear) mapping  $\Phi$  that sends vectors from the original feature space  $K$  into a higher-dimensional feature space  $K'$ .*

*Nonlinear classification in  $K$  is then obtained by creating a linear decision boundary in  $K'$ .*

*Nonlinear applications of SVM rely on an implicit (nonlinear) mapping that sends vectors from the original feature space  $K$  into a higher-dimensional feature space  $K'$ .*

*Nonlinear classification in  $K$  is then obtained by creating a linear decision boundary in  $K'$ .*

*In practice, this involves no computations in the higher dimensional space!*

# **II. MAXIMUM MARGIN HYPERPLANES**

*Q: How is the decision boundary (mmh) derived?*

*Q: How is the decision boundary (mmh) derived?*

*A: By the discriminant function,*

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b.$$

*such that  $w$  is the weight vector and  $b$  is the bias.*


*Q: How is the decision boundary (mmh) derived?*

*A: By the discriminant function,*

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b.$$

*such that  $w$  is the weight vector and  $b$  is the bias.*

*The sign of  $f(x)$  determines the (binary) class label of a record  $x$ .*



## NOTE

The weight vector determines the *orientation* of the decision boundary.

The bias determines its *translation* from the origin.



*As we said before, SVM solves for the decision boundary that minimizes generalization error, or equivalently, that has the maximum margin.*

*Q: Why are these the same thing?*

*As we said before, SVM solves for the decision boundary that minimizes generalization error, or equivalently, that has the maximum margin.*

*Q: Why are these the same thing?*

*A: Because using the mmh as the decision boundary minimizes the probability that a small perturbation in the position of a point produces a classification error.*

*As we said before, SVM solves for the decision boundary that minimizes generalization error, or equivalently, that has the maximum margin.*

### NOTE

Intuitively, the wider the margin, the clearer the distinction between classes.

*Q: Why are these the same thing?*


*A: Because using the mmh as the decision boundary minimizes the probability that a small perturbation in the position of a point produces a classification error.*

*Selecting the mmh is a straightforward exercise in analytic geometry (we won't go through the details here).*

*In particular, this task reduces to the optimization of a **convex** objective function.*

**NOTE**

The black curve  $f(x)$  is a convex function of  $x$ .



*Selecting the mmh is a straightforward exercise in analytic geometry (we won't go through the details here).*

*In particular, this task reduces to the optimization of a **convex** objective function.*

*This is nice because convex optimization problems are guaranteed to give **global optima** (and they're easy to solve numerically too).*

*Selecting the mmh is a straightforward exercise in analytic geometry (we won't go through the details here).*

*In particular, this task reduces to the optimization of a convex objective function.*


### NOTE

The heuristic techniques we've discussed (eg greedy algorithms) are not necessary with convex optimization!

*This is nice because convex optimization problems are guaranteed to give **global optima** (and they're easy to solve numerically too).*

*Notice that the margin depends only on a subset of the training data; namely, those points that are nearest to the decision boundary.*





*Notice that the margin depends only on a **subset** of the training data; namely, those points that are nearest to the decision boundary.*

*These points are called the **support vectors**.*

*Notice that the margin depends only on a **subset** of the training data; namely, those points that are nearest to the decision boundary.*

*These points are called the **support vectors**.*

*The other points (far from the decision boundary) don't affect the construction of the mmh at all!*

*All of the decision boundaries we've seen so far have split the data perfectly; eg, the data are **linearly separable**, and therefore the training error is 0.*

*All of the decision boundaries we've seen so far have split the data perfectly; eg, the data are **linearly separable**, and therefore the training error is 0.*

*The optimization problem that this SVM solves is:*

$$\begin{array}{ll} \underset{\mathbf{w}, b}{\text{minimize}} & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{subject to:} & y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \quad i = 1, \dots, n. \end{array}$$

*All of the decision boundaries we've seen so far have split the data perfectly; eg, the data are **linearly separable**, and therefore the training error is 0.*

*The optimization problem that this SVM solves is:*

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{1}{2} \|\mathbf{w}\|^2 \\ & \text{subject to:} && y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \quad i = 1, \dots, n. \end{aligned}$$

## NOTE

This type of optimization problem is called a *quadratic program*.

The result of this qp is the *hard margin classifier* we've been discussing.

# **III. SLACK VARIABLES**

*Recall that in building the hard margin classifier, we assumed that our data was **linearly separable** (eg, that we could perfectly classify each record with a linear decision boundary).*



*Recall that in building the hard margin classifier, we assumed that our data was **linearly separable** (eg, that we could perfectly classify each record with a linear decision boundary).*

*Suppose that this was not true, or suppose that we wanted to use a larger margin at the expense of incurring some training error.*

*Recall that in building the hard margin classifier, we assumed that our data was **linearly separable** (eg, that we could perfectly classify each record with a linear decision boundary).*

*Suppose that this was not true, or suppose that we wanted to use a larger margin at the expense of incurring some training error.*

*This can be done using by introducing **slack variables**.*

*Slack variables  $\xi_i$  generalize the optimization problem to permit some misclassified training records (which come at a cost  $C$ ).*

*Slack variables  $\xi_i$  generalize the optimization problem to permit some misclassified training records (which come at a cost  $C$ ).*

*The resulting soft margin classifier is given by:*


$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \\ & \text{subject to:} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0. \end{aligned}$$

*Slack variables  $\xi_i$  generalize the optimization problem to permit some misclassified training records (which come at a cost  $C$ ).*

*The resulting **soft margin classifier** is given by:*

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \\ & \text{subject to:} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0. \end{aligned}$$

*This an example of **bias-variance tradeoff**.*



*The soft-margin optimization problem can be rewritten as:*

$$\begin{array}{ll} \underset{\alpha}{\text{maximize}} & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j \\ \text{subject to:} & \sum_{i=1}^n y_i \alpha_i = 0, \quad 0 \leq \alpha_i \leq C. \end{array}$$

*The soft-margin optimization problem can be rewritten as:*

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} && \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j \\ & \text{subject to:} && \sum_{i=1}^n y_i \alpha_i = 0, \quad 0 \leq \alpha_i \leq C. \end{aligned}$$

**NOTE**

This is called the *dual formulation* of the optimization problem.

(reached via Lagrange multipliers)



*The soft-margin optimization problem can be rewritten as:*

$$\begin{array}{ll} \underset{\alpha}{\text{maximize}} & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j \\ \text{subject to:} & \sum_{i=1}^n y_i \alpha_i = 0, \quad 0 \leq \alpha_i \leq C. \end{array}$$

*Notice that this expression depends on the features  $x_i$  only via the inner product*

$$\langle x_i, x_j \rangle = x_i^T x_j$$

*The inner product is an operation that takes two vectors and returns a real number.*

*The inner product is an operation that takes two vectors and returns a real number.*

*The fact that we we can rewrite the optimization problem in terms of the inner product means that we don't actually have to do any calculations in the feature space  $K$ .*


*The inner product is an operation that takes two vectors and returns a real number.*

*The fact that we we can rewrite the optimization problem in terms of the inner product means that we don't actually have to do any calculations in the feature space  $K$ .*

*In particular, we can easily change  $K$  to be some other space  $K'$ .*

# **IV. NONLINEAR CLASSIFICATION**

*Suppose we need a more complex classifier than a linear decision boundary allows.*



*Suppose we need a more complex classifier than a linear decision boundary allows.*


*One possibility is to add nonlinear combinations of features to the data, and then to create a linear decision boundary in the enhanced (higher-dimensional) feature space.*

*Suppose we need a more complex classifier than a linear decision boundary allows.*


*One possibility is to add nonlinear combinations of features to the data, and then to create a linear decision boundary in the enhanced (higher-dimensional) feature space.*

*This linear decision boundary will be mapped to a nonlinear decision boundary in the original feature space.*





original feature space  $K$



higher-dim feature space  $K'$

*The logic of this approach is sound, but there are a few problems with this version.*

*The logic of this approach is sound, but there are a few problems with this version.*

*In particular, this will not scale well, since it requires many high-dimensional calculations.*

*Warning:*

*In particular, this will be expensive, since it requires many high-dimensional calculations.*

*It can sometimes lead to more complexity (both modeling complexity and computational complexity) than we want.*

*To recap:*

- remap the feature vectors  $x_i$  into a higher-dimensional space  $K'$*
- create a linear decision boundary in  $K'$*
- back out the nonlinear decision boundary in  $K$  from the result*

*Recall that our optimization problem depends on the features only through the inner product  $x^T x$ :*

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} && \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j \\ & \text{subject to:} && \sum_{i=1}^n y_i \alpha_i = 0, \quad 0 \leq \alpha_i \leq C. \end{aligned}$$

*Recall that our optimization problem depends on the features only through the inner product  $x^T x$ :*

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} && \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j \\ & \text{subject to:} && \sum_{i=1}^n y_i \alpha_i = 0, \quad 0 \leq \alpha_i \leq C. \end{aligned}$$

*We can replace this inner product with a more general function that has the same type of output as the inner product.*

*Formally, we can think of the inner product as a map that sends two vectors in the feature space  $K$  into the real line  $\mathbb{R}$ .*



*Formally, we can think of the inner product as a map that sends two vectors in the feature space  $K$  into the real line  $\mathbb{R}$ .*

*We can replace this with a generalization of the inner product called a **kernel function** that maps two vectors in a higher-dimensional feature space  $K'$  into  $\mathbb{R}$ .*

*The upshot is that we can use a kernel function to implicitly train our model in a higher-dimensional feature space, without incurring additional computational complexity!*

*The upshot is that we can use a kernel function to *implicitly* train our model in a higher-dimensional feature space, *without* incurring additional computational complexity!*

*As long as the kernel function satisfies certain conditions, our conclusions above regarding the mmh continue to hold.*

*The upshot is that we can use a kernel function to **implicitly** train our model in a higher-dimensional feature space, **without** incurring additional computational complexity!*

*As long as the kernel function satisfies certain conditions, our conclusions above regarding the mmh continue to hold.*

### NOTE

These conditions are contained in a result called *Mercer's theorem*.

*The upshot is that we can use a kernel function to implicitly train our model in a higher-dimensional feature space, without incurring additional computational complexity!*

*As long as the kernel function satisfies certain conditions, our conclusions above regarding the mmh continue to hold.*

*In other words, no algorithmic changes are necessary, and all the benefits of a linear SVM are maintained.*

*some popular kernels:*

*linear kernel*                       $k(\mathbf{x}, \mathbf{x}') = \langle \mathbf{x}, \mathbf{x}' \rangle$

*polynomial kernel*               $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + 1)^d$

*Gaussian kernel*                 $k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$


*some popular kernels:*

*linear kernel*                       $k(\mathbf{x}, \mathbf{x}') = \langle \mathbf{x}, \mathbf{x}' \rangle$





*polynomial kernel*               $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + 1)^d$

*Gaussian kernel*                 $k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$

*The hyperparameters  $d, \gamma$  affect the flexibility of the decision bdy.*







*SVMs (and **kernel methods** in general) are versatile, powerful, and popular techniques that can produce accurate results for a wide array of classification problems.*

*The main disadvantage of SVMs is the lack of intuition they produce. These models are truly black boxes!*

---

**INTRO TO DATA SCIENCE**

---

**EX: SVM IN SCIKIT-LEARN**